

Dimensional Modeling Report – Instamart (Snowflake + S3)

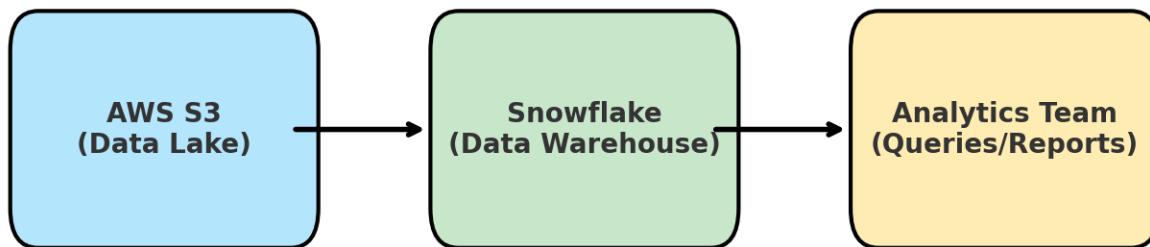
Introduction

Instacart is a grocery ordering and delivery app that helps customers conveniently stock their refrigerators and pantries with their favourite items and daily essentials. After customers place orders through the Instacart app, personal shoppers pick the products in-store and deliver them directly to the customer's doorstep.

To support the Instamart analytics team, we need to design a dimensional model for a data warehouse. This will allow the team to run their analytics efficiently, ensuring smooth performance and faster query execution. An optimized schema is essential to provide insights at minimal execution time.

For this project, we will be leveraging **Snowflake** as our cloud data warehouse solution, with raw data stored in an **AWS S3 data lake**. This setup will enable scalable storage, seamless integration, and high-performance analytics.

Data Flow: S3 → Snowflake → Analytics



1. Business Requirements

Using this business needs answers to the following questions from analytics team and to design a Tableau dashboard if needed to answer the following questions

Business Questions for Instamart Dimensional Model

Customer Behaviour & Ordering Patterns

- What are the most frequently reordered products?
- How often do customers place repeat orders (days between orders)?
- At what time of day and day of week do customers place the most orders?
- What is the average number of items in a customer's cart per order?

Product & Category Insights

- Which products are most commonly added to carts first?
- Which **aisles** and **departments** generate the highest order volumes?
- What are the top-selling products by number of orders?
- Are certain product categories more likely to be reordered?

Operational & Inventory Questions

- How many products are ordered per department/aisle in a given week or month?
- What percentage of products in each department are reordered at least once?
- Which aisles contribute most to overall sales volume?

Trend & Performance Tracking

- What is the trend of total orders over time (daily, weekly, monthly)?
- How does order frequency vary by season, day, or hour?
- What percentage of orders contain products from multiple departments?

Key Metrics (Facts)

The fact table will mainly be built from orders and order_products since they capture measurable events.

- Number of Orders → Count of order_id
- Number of Products per Order → Count of product_id grouped by order_id
- Cart Position → From add_to_cart_order (sequence of product placement in cart)
- Reorder Rate → Based on reordered flag (0 = new order, 1 = reordered)
- Average Days Between Orders → From days_since_prior_order
- Orders by Time/Day → Derived from order_dow and order_hour_of_day

(Note: Since no pricing/revenue data exists in these tables, facts are focused on order frequency, product mix, and reorders.)

3. Key Categories (Dimensions)

Customer Dimension (derived from orders)

- user_id → Unique customer
- order_number → Sequence of orders per customer
- Behavior attributes can be derived (e.g., frequency, reorder tendency).

Product Dimension (*from products, linked to aisles & departments*)

- product_id
- product_name
- aisle_id → joined with Aisles table → aisle (e.g., beverages, dairy)
- department_id → joined with Departments table → department (e.g., produce, frozen)

Time Dimension (*from orders*)

- order_dow → Day of week (0 = Sunday, 6 = Saturday)
- order_hour_of_day → Hour product was ordered (0–23)
- Derived hierarchy: Day → Week → Month → Year

Order Dimension (*from orders & order_products*)

- order_id
- eval_set (train, test, prior – depending on dataset split)
- add_to_cart_order (cart sequence)
- reordered (yes/no)

Aisle Dimension

- aisle_id
- aisle (e.g., snacks, dairy, frozen foods)

Department Dimension

- department_id
- department (e.g., produce, household, beverages)

Fact Tables & Measures

1) FactOrderItems

Grain: one record per (order_id, product_id)

Keys

- order_id → (joins to Order/Time/Customer dimensions)
- product_id → (joins to Product, Aisle, Department dims)

Measures

- item_count (always 1 per row; useful for summing quantity)
- add_to_cart_order (sequence position)
- reordered (0/1; supports reorder rate)

Degenerate attributes (optional)

- None required (kept in Order dim/degenerate order)

Common derived KPIs

- Reorder rate = avg(reordered)
- Average cart position = avg(add_to_cart_order)
- Items per order = sum(item_count) by order_id

2. Dimension Tables

ProductDim

From **products** (enriched with aisle/department names)

- Keys: product_id (surrogate product_key recommended)
- Attributes: product_name, aisle_id, aisle_name, department_id, department_name
- Hierarchies: **Department → Aisle → Product**

AisleDim (optional if not denormalizing into ProductDim)

- Keys: aisle_id (surrogate aisle_key)
- Attributes: aisle

DepartmentDim (optional if not denormalizing into ProductDim)

- Keys: department_id (surrogate department_key)
- Attributes: department

In a star schema you can **denormalize** aisle/department into ProductDim to keep joins simple; keep separate dims only if they're reused or have their own attributes.

TimeDim

Derived from order timestamps

- Keys: date_key (YYYYMMDD), optionally time_key (hour)
- Attributes: date, day_name, order_dow, week, month, quarter, year, order_hour_of_day, is_weekend

CustomerDim (*if/when you have user attributes*)

- Keys: user_id (surrogate customer_key)
- Attributes: e.g., signup_date, segment (if available later)
- Until you have user profile data, keep it minimal: just user_id

OrderDegenerateDim (very light)

- Degenerate attributes that don't merit a full dimension:
 - order_id (often treated as a **degenerate dimension** on facts)
 - order_number, eval_set (if present)

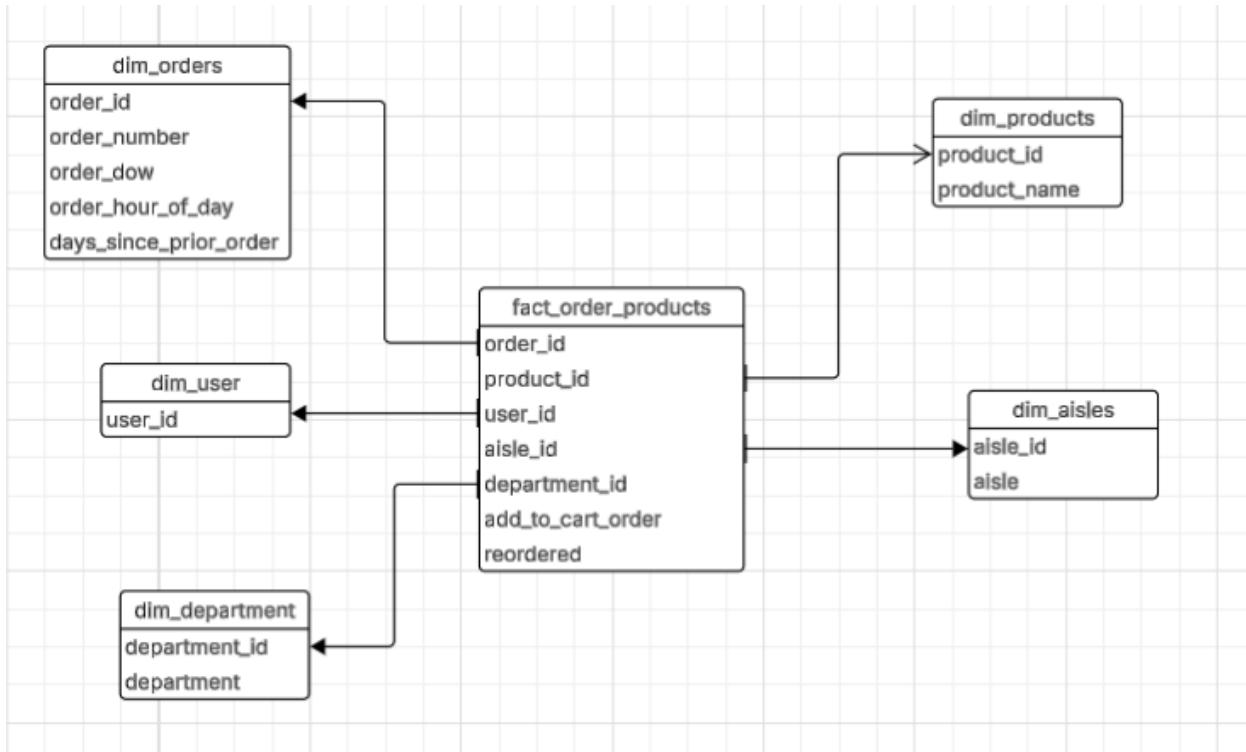
Notes for Snowflake Implementation

- Staging: Load CSVs from S3 stage → Snowflake staging tables.
- ELT: Build dims first (ProductDim; optionally Aisle/Department dims; TimeDim), then FactOrderItems, then FactOrders.

- Surrogate keys: Use IDENTITY or SEQUENCE for dimension *_key columns; retain natural keys (product_id, etc.) for lineage.
- Performance: Consider clustering on (order_date_key, product_key) or (order_date_key, department_name) for heavy time/category queries.
- SCDs: If product/category attributes change, start with Type 1 (overwrite) and evolve to Type 2 if history is needed.

3. Schema Diagram

Star Schema:



We are using a **Star Schema** for the Instamart data warehouse because it is the most effective design for analytical workloads such as reporting, dashboards, and ad-hoc queries.

1. Query Performance

- Star schemas reduce the number of joins compared to normalized schemas.
- Facts are stored in a central Fact table, while descriptive attributes are in smaller Dimension tables.
- This structure allows Snowflake's query engine to scan less data, improving speed for common analytical queries (e.g., sales by aisle, orders by hour).

2. Simplicity & Usability

- Easy for analysts and business users to understand: one central fact connected to multiple descriptive dimensions.
- Intuitive structure: “What happened?” (facts) vs. “How/Who/When/Where?” (dimensions).

3. Flexibility in Analysis

- Supports slice-and-dice analysis across different dimensions (e.g., orders by product, by department, by time).
- Enables drill-down hierarchies (e.g., Department → Aisle → Product, or Day → Month → Year).

4. Scalability with Snowflake

- Star schemas are efficient in columnar cloud warehouses like Snowflake, which handle large fact tables with billions of rows.
- Dimensions are relatively small and can be cached, minimizing repeated scans.

5. Extensibility

- Easy to add new dimensions (e.g., Customer profile, Delivery performance) or new facts (e.g., Payment transactions) without redesigning the entire schema.

Conclusion

The dimensional model for Instamart is designed using a **Star Schema** with a central fact table (orders and order line items) surrounded by descriptive dimensions (products, time, customer, aisle, department, and order details). This design ensures that the Instamart analytics team can perform queries quickly, explore data from multiple perspectives, and generate actionable insights with minimal execution time.

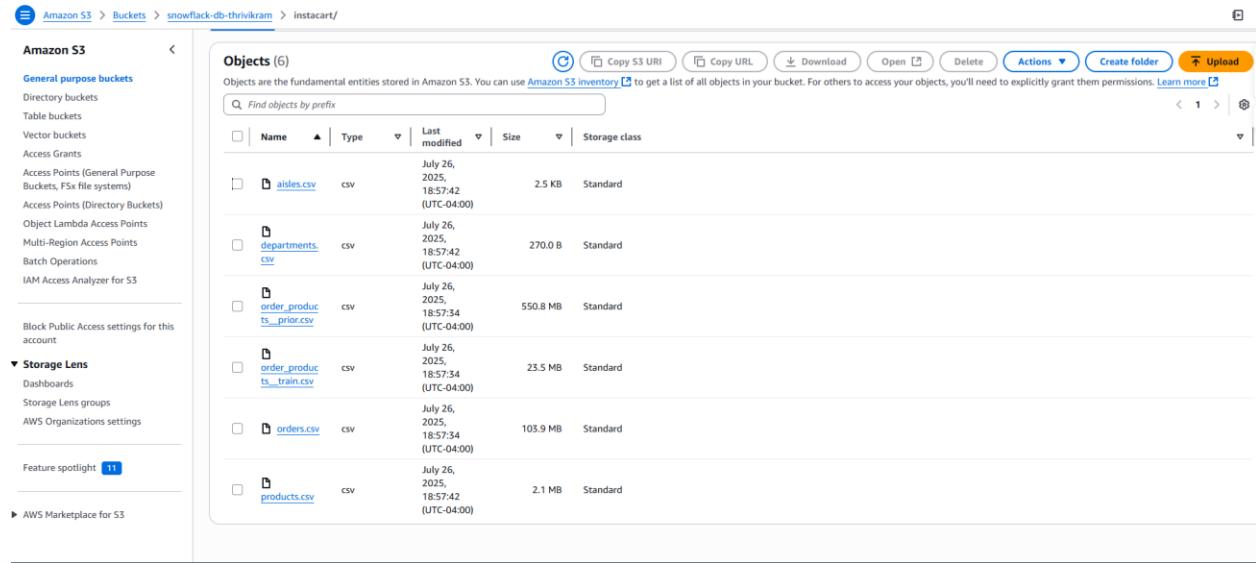
By leveraging **Snowflake** as the data warehouse and **AWS S3** as the data lake, we achieve a scalable, cloud-native solution that supports both large data volumes and flexible analytics. The star schema structure also provides simplicity, extensibility, and strong performance — making it well-suited for ongoing reporting and research needs.

The following **Evidence & Appendix** section will include screenshots of the implementation, SQL scripts for building fact and dimension tables, and example analytical queries to demonstrate the effectiveness of this model.

Evidence & SQL Appendix (Process Walkthrough)

This section documents the step-by-step process we followed to build the dimensional model in Snowflake using data staged in AWS S3. Each step includes an explanation, the SQL code used, and placeholders for screenshots (SS) as evidence.

Below is a SS of files in S3

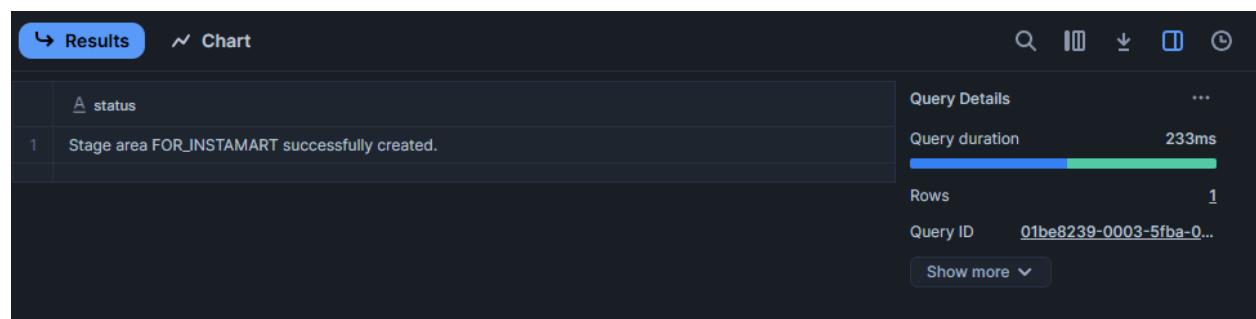


Name	Type	Last modified	Size	Storage class
aisles.csv	csv	July 26, 2025, 18:57:42 (UTC-04:00)	2.5 KB	Standard
departments.csv	csv	July 26, 2025, 18:57:42 (UTC-04:00)	270.0 B	Standard
order_products_prior.csv	csv	July 26, 2025, 18:57:34 (UTC-04:00)	550.8 MB	Standard
order_products_train.csv	csv	July 26, 2025, 18:57:34 (UTC-04:00)	23.5 MB	Standard
orders.csv	csv	July 26, 2025, 18:57:34 (UTC-04:00)	103.9 MB	Standard
products.csv	csv	July 26, 2025, 18:57:42 (UTC-04:00)	2.1 MB	Standard

Step 1: Stage Creation

We first created a Snowflake stage pointing to our S3 bucket, which serves as the raw data lake.

```
CREATE OR REPLACE STAGE for_instamart
URL = 's3://snowflake-db-thrividram/instacart/'
CREDENTIALS =
  AWS_KEY_ID = '<REDACTED>',
  AWS_SECRET_KEY = '<REDACTED>'
);
```



Step 2: File Format Definition

Defined a CSV file format to consistently parse incoming files.

```
CREATE OR REPLACE FILE FORMAT insta_cart_csv
  TYPE = 'CSV'
  FIELD_DELIMITER = ','
  SKIP_HEADER = 1
  FIELD_OPTIONALLY_ENCLOSED_BY = '\"';
```

The screenshot shows a database query results interface. At the top, there are tabs for 'Results' and 'Chart'. Below the tabs, the results table has a single row with the following data:

	status
1	File format INSTA_CART_CSV successfully created.

On the right side of the results table, there is a 'Query Details' panel with the following information:

- Query duration: 83ms
- Rows: 1
- Query ID: 01be8242-0003-5fb9-0...

At the bottom right of the results table, there is a 'Show more' button.

Step 3: Base Table Creation

Created raw tables to hold aisles, departments, products, orders, and order-product mappings.

```
CREATE OR REPLACE TABLE aisles (
  aisle_id INTEGER,
  aisle   VARCHAR
);

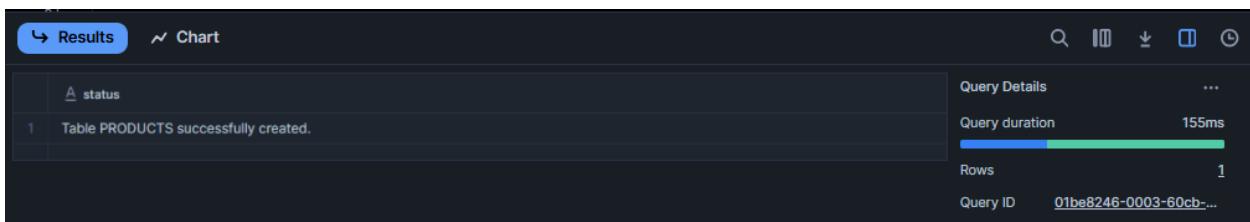
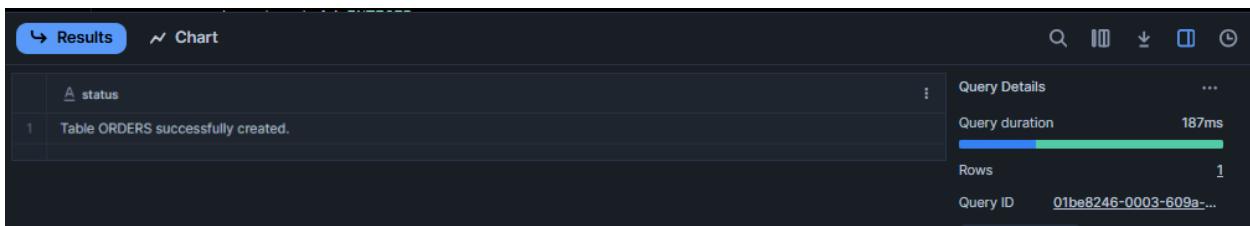
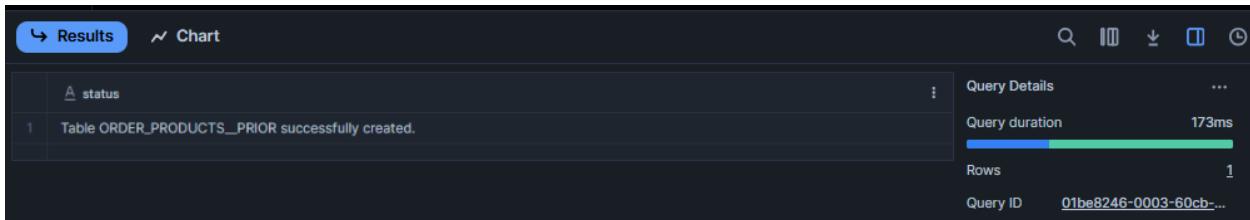
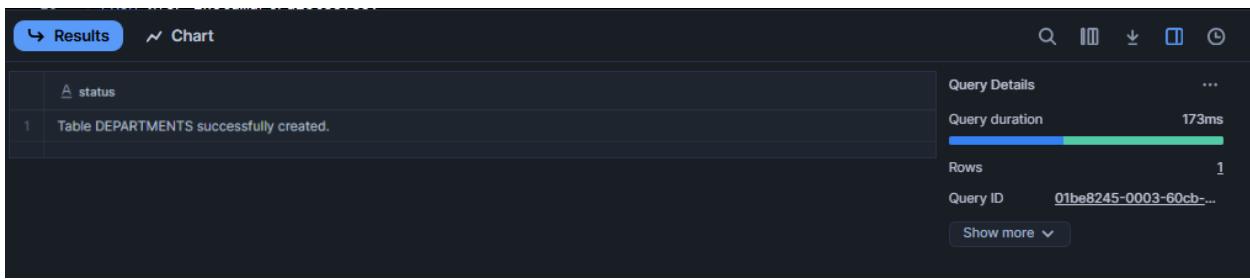
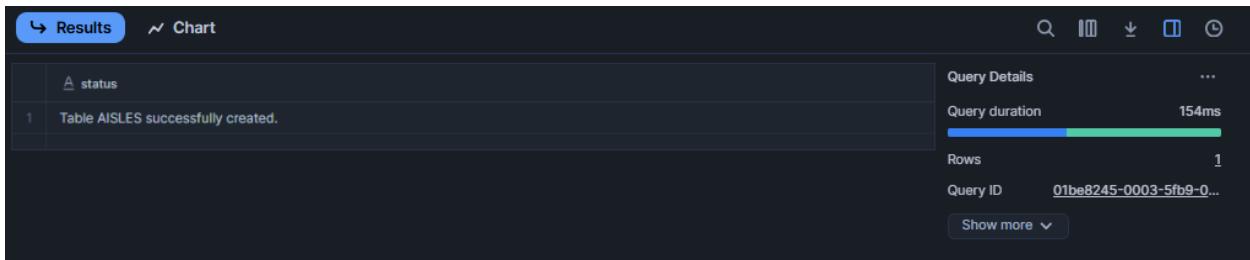
CREATE OR REPLACE TABLE departments (
  department_id INTEGER,
  department   VARCHAR
);

CREATE OR REPLACE TABLE products (
  product_id    INTEGER,
  product_name  VARCHAR,
  aisle_id      INTEGER,
  department_id INTEGER
);

CREATE OR REPLACE TABLE orders (
  order_id        INTEGER,
  user_id         INTEGER,
  eval_set        VARCHAR,
  order_number    INTEGER,
  order_dow       INTEGER,
  order_hour_of_day INTEGER,
  days_since_prior_order INTEGER
);

CREATE OR REPLACE TABLE order_products_prior (
  order_id        INTEGER,
  product_id     INTEGER,
  add_to_cart_order INTEGER,
  reordered      INTEGER
);
```

Results:



Step 4: Loading Data from S3

We loaded CSV files from S3 into Snowflake tables using COPY INTO.

```
COPY INTO aisles (aisle_id, aisle)
FROM @for_instamart/aisles.csv
FILE_FORMAT = (FORMAT_NAME = 'insta_cart_csv');

COPY INTO departments
FROM @for_instamart/departments.csv
FILE_FORMAT = (FORMAT_NAME = 'insta_cart_csv');

COPY INTO products (product_id, product_name, aisle_id, department_id)
FROM @for_instamart/products.csv
FILE_FORMAT = (FORMAT_NAME = 'insta_cart_csv');

COPY INTO orders
FROM @for_instamart/orders.csv
FILE_FORMAT = (FORMAT_NAME = 'insta_cart_csv');

COPY INTO order_products_prior
FROM @for_instamart/order_products_prior.csv
FILE_FORMAT = (FORMAT_NAME = 'insta_cart_csv');
```

Results:

The figure consists of three vertically stacked screenshots of the Snowflake UI, each showing the results of a COPY INTO operation. Each screenshot has a header bar with 'Results' and 'Chart' tabs, search, filter, and refresh icons.

Screenshot 1: aisles

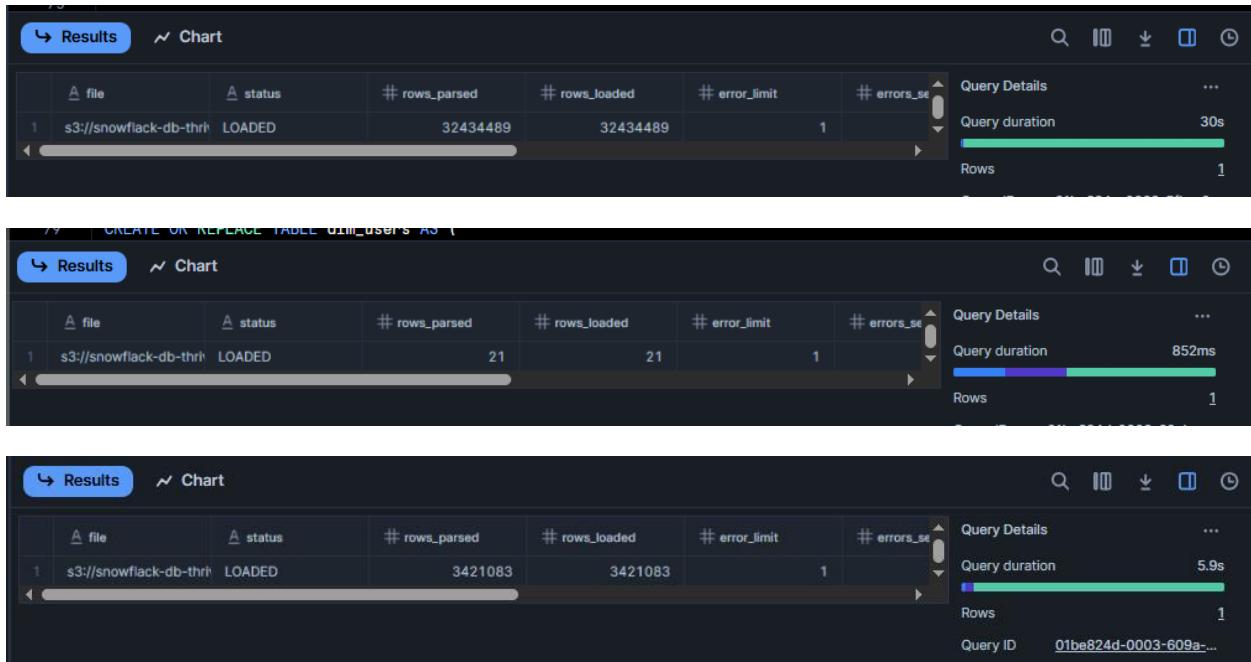
	file	status	# rows_parsed	# rows_loaded	# error_limit	# errors_seen	Query Details	...
1	s3://snowflake-db-thr\	LOADED	134	134	1	0	Query duration	975ms

Screenshot 2: departments

	file	status	# rows_parsed	# rows_loaded	# error_limit	# errors_seen	Query Details	...
1	s3://snowflake-db-thr\	LOADED	49688	49688	1	0	Query duration	1.1s

Screenshot 3: products

	file	status	# rows_parsed	# rows_loaded	# error_limit	# errors_seen	Query Details	...
1	s3://snowflake-db-thr\	LOADED	1384617	1384617	1	0	Query duration	2.6s



79

CREATE OR REPLACE TABLE dim_users AS (

	file	status	# rows_parsed	# rows_loaded	# error_limit	# errors_seen
1	s3://snowflake-db-thr1	LOADED	32434489	32434489	1	0

Query Details ...
Query duration 30s
Rows 1

CREATE OR REPLACE TABLE dim_products AS (

	file	status	# rows_parsed	# rows_loaded	# error_limit	# errors_seen
1	s3://snowflake-db-thr1	LOADED	21	21	1	0

Query Details ...
Query duration 852ms
Rows 1

CREATE OR REPLACE TABLE dim_aisles AS (

	file	status	# rows_parsed	# rows_loaded	# error_limit	# errors_seen
1	s3://snowflake-db-thr1	LOADED	3421083	3421083	1	0

Query Details ...
Query duration 5.9s
Rows 1
Query ID 01be824d-0003-609a-...

Step 5: Building Dimensions

We created dimension tables from the base tables.

```
CREATE OR REPLACE TABLE dim_users AS (
    SELECT DISTINCT user_id FROM orders
);

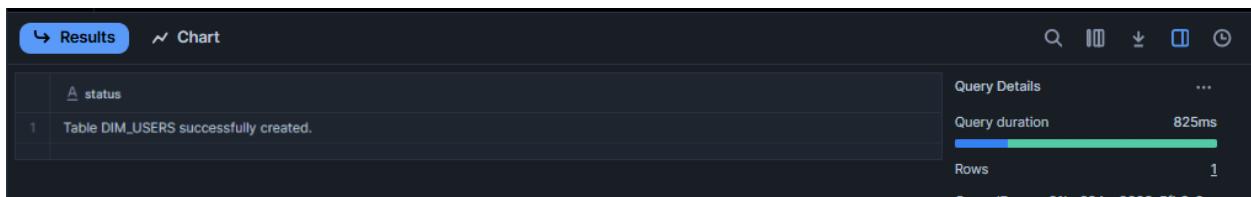
CREATE OR REPLACE TABLE dim_products AS (
    SELECT product_id, product_name FROM products
);

CREATE OR REPLACE TABLE dim_aisles AS (
    SELECT aisle_id, aisle FROM aisles
);

CREATE OR REPLACE TABLE dim_departments AS (
    SELECT department_id, department FROM departments
);

CREATE OR REPLACE TABLE dim_orders AS (
    SELECT order_id, order_number, order_dow, order_hour_of_day, days_since_prior_order
    FROM orders
);
```

Results:



	status
1	Table DIM_USERS successfully created.

Query Details ...
Query duration 825ms
Rows 1
Query ID 01be824d-0003-609a-...

Results Chart

status
1 Table DIM_PRODUCTS successfully created.

Query Details

Query duration 738ms

Rows 1

Results Chart

status
1 Table DIM_AISLES successfully created.

Query Details

Query duration 747ms

Rows 1

Results Chart

status
1 Table DIM_DEPARTMENTS successfully created.

Query Details

Query duration 574ms

Rows 1

Results Chart

status
1 Table DIM_ORDERS successfully created.

Query Details

Query duration 778ms

Rows 1

Step 6: Building the Fact Table

We joined orders, products, and order-product mappings to construct the fact table.

```
CREATE OR REPLACE TABLE fact_order_product AS (
  SELECT
    op.order_id,
    op.product_id,
    o.user_id,
    p.department_id,
    p.aisle_id,
    op.add_to_cart_order,
    op.reordered
  FROM order_products_prior op
  JOIN orders o ON op.order_id = o.order_id
  JOIN products p ON op.product_id = p.product_id
);
```

Results		Chart	Query Details
	status	...	Query duration
1	Table FACT_ORDER_PRODUCT successfully created.		5.0s

Step 7: QA Checks & Sample Analytics

```
-- Count checks
SELECT 'orders' AS tbl, COUNT(*) AS count FROM orders
UNION ALL SELECT 'products', COUNT(*) FROM products
UNION ALL SELECT 'order_products_prior', COUNT(*) FROM order_products_prior;

-- Distinct keys
SELECT COUNT(DISTINCT product_id) AS distinct_products FROM products;
SELECT COUNT(DISTINCT order_id) AS distinct_orders FROM orders;

-- Example: Reorder rate by product
SELECT p.product_name, AVG(op.reordered) AS reorder_rate
FROM fact_order_product op
JOIN dim_products p USING (product_id)
GROUP BY 1
ORDER BY 2 DESC
LIMIT 20;

-- Example: Peak ordering hour
SELECT d.order_hour_of_day, COUNT(*) AS orders_at_hour
FROM dim_orders d
GROUP BY 1
ORDER BY 2 DESC;
```

Results		Chart	Query Details
	# TBL	Count(*)	...
1	orders	3421083	Query duration
2	products	49688	297ms
3	order_products_prior	33819106	Rows
			Query ID

Results		Chart	Query Details
	# DISTINCT_ORDERS	Count(*)	...
1		3421083	Query duration
			231ms
			Rows
			1

Results Chart

Query Details

- Query duration 1.1s
- Rows 20
- Query ID 01be8252-0003-612f-0...

PRODUCT_NAME	REORDER_RATE
1 Raw Veggie Wrappers	0.942029
2 Serenity Ultimate Extrema Overnight Pads	0.933333
3 Orange Energy Shots	0.923077
4 Chocolate Love Bar	0.921569

Results Chart

Query Details

- Query duration 166ms
- Rows 24
- Query ID 01be8252-0003-612f-0...

ORDER_HOUR_OF_DAY	ORDERS_AT_HOUR
10	288418
11	284728
15	283639
14	283042

References

- Kimball & Ross, The Data Warehouse Toolkit
- Snowflake Docs: Loading data from S3, Clustering, SCD patterns